

Solver.py

Untitled1

February 3, 2021

```
[ ]: from __future__ import print_function, division
from builtins import range
from builtins import object
import os
import pickle as pickle

import numpy as np

from nnrl import optim

class Solver(object):
    """
    A Solver encapsulates all the logic necessary for training classification
    models. The Solver performs stochastic gradient descent using different
    update rules defined in optim.py.

    The solver accepts both training and validation data and labels so it can
    periodically check classification accuracy on both training and validation
    data to watch out for overfitting.

    To train a model, you will first construct a Solver instance, passing the
    model, dataset, and various options (learning rate, batch size, etc) to the
    constructor. You will then call the train() method to run the optimization
    procedure and train the model.

    After the train() method returns, model.params will contain the parameters
    that performed best on the validation set over the course of training.
    In addition, the instance variable solver.loss_history will contain a list
    of all losses encountered during training and the instance variables
    solver.train_acc_history and solver.val_acc_history will be lists of the
    accuracies of the model on the training and validation set at each epoch.

    Example usage might look something like this:

    data = {
        'X_train': # training data
```

```

    'y_train': # training labels
    'X_val': # validation data
    'y_val': # validation labels
}
model = MyAwesomeModel(hidden_size=100, reg=10)
solver = Solver(model, data,
                 update_rule='sgd',
                 optim_config={
                     'learning_rate': 1e-3,
                 },
                 lr_decay=0.95,
                 num_epochs=10, batch_size=100,
                 print_every=100)
solver.train()

```

A Solver works on a model object that must conform to the following API:

- `model.params` must be a dictionary mapping string parameter names to numpy arrays containing parameter values.
- `model.loss(X, y)` must be a function that computes training-time loss and gradients, and test-time classification scores, with the following inputs and outputs:

Inputs:

- `X`: Array giving a minibatch of input data of shape (N, d_1, \dots, d_k)
- `y`: Array of labels, of shape $(N,)$ giving labels for `X` where `y[i]` is the label for `X[i]`.

Returns:

If `y` is `None`, run a test-time forward pass and return:

- `scores`: Array of shape (N, C) giving classification scores for `X` where `scores[i, c]` gives the score of class `c` for `X[i]`.

If `y` is not `None`, run a training time forward and backward pass and return a tuple of:

- `loss`: Scalar giving the loss
- `grads`: Dictionary with the same keys as `self.params` mapping parameter names to gradients of the loss with respect to those parameters.

"""

```

def __init__(self, model, data, **kwargs):
    """

```

Construct a new Solver instance.

Required arguments:

- *model*: A model object conforming to the API described above
- *data*: A dictionary of training and validation data containing:
 - 'X_train': Array, shape (N_train, d_1, ..., d_k) of training images
 - 'X_val': Array, shape (N_val, d_1, ..., d_k) of validation images
 - 'y_train': Array, shape (N_train,) of labels for training images
 - 'y_val': Array, shape (N_val,) of labels for validation images

Optional arguments:

- *update_rule*: A string giving the name of an update rule in *optim.py*. Default is 'sgd'.
- *optim_config*: A dictionary containing hyperparameters that will be passed to the chosen update rule. Each update rule requires different hyperparameters (see *optim.py*) but all update rules require a 'learning_rate' parameter so that should always be present.
- *lr_decay*: A scalar for learning rate decay; after each epoch the learning rate is multiplied by this value.
- *batch_size*: Size of minibatches used to compute loss and gradient during training.
- *num_epochs*: The number of epochs to run for during training.
- *print_every*: Integer; training losses will be printed every *print_every* iterations.
- *verbose*: Boolean; if set to false then no output will be printed during training.
- *num_train_samples*: Number of training samples used to check training accuracy; default is 1000; set to None to use entire training set.
- *num_val_samples*: Number of validation samples to use to check val accuracy; default is None, which uses the entire validation set.
- *checkpoint_name*: If not None, then save model checkpoints here every epoch.

"""

```

self.model = model
self.X_train = data['X_train']
self.y_train = data['y_train']
self.X_val = data['X_val']
self.y_val = data['y_val']

# Unpack keyword arguments
self.update_rule = kwargs.pop('update_rule', 'sgd')
self.optim_config = kwargs.pop('optim_config', {})
self.lr_decay = kwargs.pop('lr_decay', 1.0)
self.batch_size = kwargs.pop('batch_size', 100)
self.num_epochs = kwargs.pop('num_epochs', 10)
self.num_train_samples = kwargs.pop('num_train_samples', 1000)
self.num_val_samples = kwargs.pop('num_val_samples', None)

self.checkpoint_name = kwargs.pop('checkpoint_name', None)
self.print_every = kwargs.pop('print_every', 10)

```

```

self.verbose = kwargs.pop('verbose', True)

# Throw an error if there are extra keyword arguments
if len(kwargs) > 0:
    extra = ', '.join("%s" % k for k in list(kwargs.keys()))
    raise ValueError('Unrecognized arguments %s' % extra)

# Make sure the update rule exists, then replace the string
# name with the actual function
if not hasattr(optim, self.update_rule):
    raise ValueError('Invalid update_rule "%s" % self.update_rule)
self.update_rule = getattr(optim, self.update_rule)

self._reset()

def _reset(self):
    """
    Set up some book-keeping variables for optimization. Don't call this
    manually.
    """
    # Set up some variables for book-keeping
    self.epoch = 0
    self.best_val_acc = 0
    self.best_params = {}
    self.loss_history = []
    self.train_acc_history = []
    self.val_acc_history = []

    # Make a deep copy of the optim_config for each parameter
    self.optim_configs = {}
    for p in self.model.params:
        d = {k: v for k, v in self.optim_config.items()}
        self.optim_configs[p] = d

def _step(self):
    """
    Make a single gradient update. This is called by train() and should not
    be called manually.
    """
    # Make a minibatch of training data
    num_train = self.X_train.shape[0]
    batch_mask = np.random.choice(num_train, self.batch_size)
    X_batch = self.X_train[batch_mask]
    y_batch = self.y_train[batch_mask]

```

```

    # Compute loss and gradient
    loss, grads = self.model.loss(X_batch, y_batch)
    self.loss_history.append(loss)

    # Perform a parameter update
    for p, w in self.model.params.items():
        dw = grads[p]
        config = self.optim_configs[p]
        w = np.reshape(w, np.shape(dw) )           # WATCH OUT HERE !!!!!!!
        #print('shapes of w, dw, config are: ', np.shape(w), np.shape(dw),
        ↪ np.shape(config['learning_rate'])) )
        next_w, next_config = self.update_rule(w, dw, config)
        self.model.params[p] = next_w
        self.optim_configs[p] = next_config

def _save_checkpoint(self):
    if self.checkpoint_name is None: return
    checkpoint = {
        'model': self.model,
        'update_rule': self.update_rule,
        'lr_decay': self.lr_decay,
        'optim_config': self.optim_config,
        'batch_size': self.batch_size,
        'num_train_samples': self.num_train_samples,
        'num_val_samples': self.num_val_samples,
        'epoch': self.epoch,
        'loss_history': self.loss_history,
        'train_acc_history': self.train_acc_history,
        'val_acc_history': self.val_acc_history,
    }
    filename = '%s_epoch%d.pkl' % (self.checkpoint_name, self.epoch)
    if self.verbose:
        print('Saving checkpoint to "%s"' % filename)
    with open(filename, 'wb') as f:
        pickle.dump(checkpoint, f)

def check_accuracy(self, X, y, num_samples=None, batch_size=100):
    """
    Check accuracy of the model on the provided data.

    Inputs:
    - X: Array of data, of shape (N, d_1, ..., d_k)
    - y: Array of labels, of shape (N,)
    - num_samples: If not None, subsample the data and only test the model
      on num_samples datapoints.

```

```

- batch_size: Split X and y into batches of this size to avoid using
  too much memory.

Returns:
- acc: Scalar giving the fraction of instances that were correctly
  classified by the model.
"""

# Maybe subsample the data
N = X.shape[0]
if num_samples is not None and N > num_samples:
    mask = np.random.choice(N, num_samples)
    N = num_samples
    X = X[mask]
    y = y[mask]

# Compute predictions in batches
num_batches = N // batch_size
if N % batch_size != 0:
    num_batches += 1
y_pred = []
for i in range(num_batches):
    start = i * batch_size
    end = (i + 1) * batch_size
    scores = self.model.loss(X[start:end])
    y_pred.append(np.argmax(scores, axis=1))
y_pred = np.hstack(y_pred)
acc = np.mean(y_pred == y)

return acc

def train(self):
    """
    Run optimization to train the model.
    """
    num_train = self.X_train.shape[0]
    iterations_per_epoch = max(num_train // self.batch_size, 1)
    num_iterations = self.num_epochs * iterations_per_epoch

    for t in range(num_iterations):
        self._step()

        # Maybe print training loss
        if self.verbose and t % self.print_every == 0:
            print('(Iteration %d / %d) loss: %f' % (
                t + 1, num_iterations, self.loss_history[-1]))

```

```

# At the end of every epoch, increment the epoch counter and decay
# the learning rate.
epoch_end = (t + 1) % iterations_per_epoch == 0
if epoch_end:
    self.epoch += 1
    for k in self.optim_configs:
        self.optim_configs[k]['learning_rate'] *= self.lr_decay

# Check train and val accuracy on the first iteration, the last
# iteration, and at the end of each epoch.
first_it = (t == 0)
last_it = (t == num_iterations - 1)
if first_it or last_it or epoch_end:
    train_acc = self.check_accuracy(self.X_train, self.y_train,
                                    num_samples=self.num_train_samples)
    val_acc = self.check_accuracy(self.X_val, self.y_val,
                                  num_samples=self.num_val_samples)
    self.train_acc_history.append(train_acc)
    self.val_acc_history.append(val_acc)
    self._save_checkpoint()

    if self.verbose:
        print('(Epoch %d / %d) train acc: %f; val_acc: %f' % (
            self.epoch, self.num_epochs, train_acc, val_acc))

# Keep track of the best model
if val_acc > self.best_val_acc:
    self.best_val_acc = val_acc
    self.best_params = {}
    for k, v in self.model.params.items():
        self.best_params[k] = v.copy()

# At the end of training swap the best params into the model
self.model.params = self.best_params

```